



Scalable Parallel Suffix Array Construction

Fabian Kulla and Peter Sanders
Karlsruhe Institute of Technology

Karlsruhe Institute of Technology



Outline



- Suffix array
- Motivation
- Parallel Difference Cover 3 algorithm
- Evaluation

Karlsruhe Institute of Technology

Suffix array (example)

Index	String	Suffix array	sorted Suffixes
0	B	5	A
1	A	3	A N A
2	N	1	A N A N A
3	A	0	B A N A N A
4	N	4	N A
5	A	2	N A N A



Motivation



- Fast phrase search in text databases in logarithmical time
 - (Find “to be or not to be”)
- Data compression
 - (e.g. Burrows Wheeler transform)
- Bioinformatics
 - Genome analysis



Sample sort (deterministic)



Input uniformly distributed

1. Sort local
2. Select samples
3. Collect and sort samples
4. Get / Broadcast splitters
5. Partition with splitters
6. Collective exchange of buckets
all-to-all
7. Merge buckets

P0	P1	P2
47 11 0 8 15 9	11 3 50 45 38 60	24 12 34 53 37 28
0 8 9 11 15 47	3 11 38 45 50 60	12 24 28 34 37 53
0 3 11 12 34 45		
0 8 9 11 15 47	3 11 38 45 50 60	12 24 28 34 37 53
0 8 9 11 3 11	15 38 45 12 24 28 34 37	47 50 60 53
0 3 8 9 11 11	12 15 24 28 34 37 38 45	47 50 53 60



Difference Cover 3 Algorithm (DC3)



- Sort suffixes starting at position $i \bmod 3 \neq 0$
- Sort remaining suffixes with help of already sorted suffix
- Merge both sorted suffix groups



Difference Cover 3 Algorithm (DC3)

- Sort suffixes starting at position $i \bmod 3 \neq 0$

- Sorting triples, recursion with 2/3 of input

0 1 2 3 4 5
B A N A N A

mod1-suffix: ANANA , NA

mod2-suffix: NANA , A

- Sort remaining suffixes with help of already sorted suffix

- Sorting pairs

B A N A N A

2 4 3 1

$(A,3) < (B,2)$

- Merge both sorted suffix groups

- Comparing pairs and triples

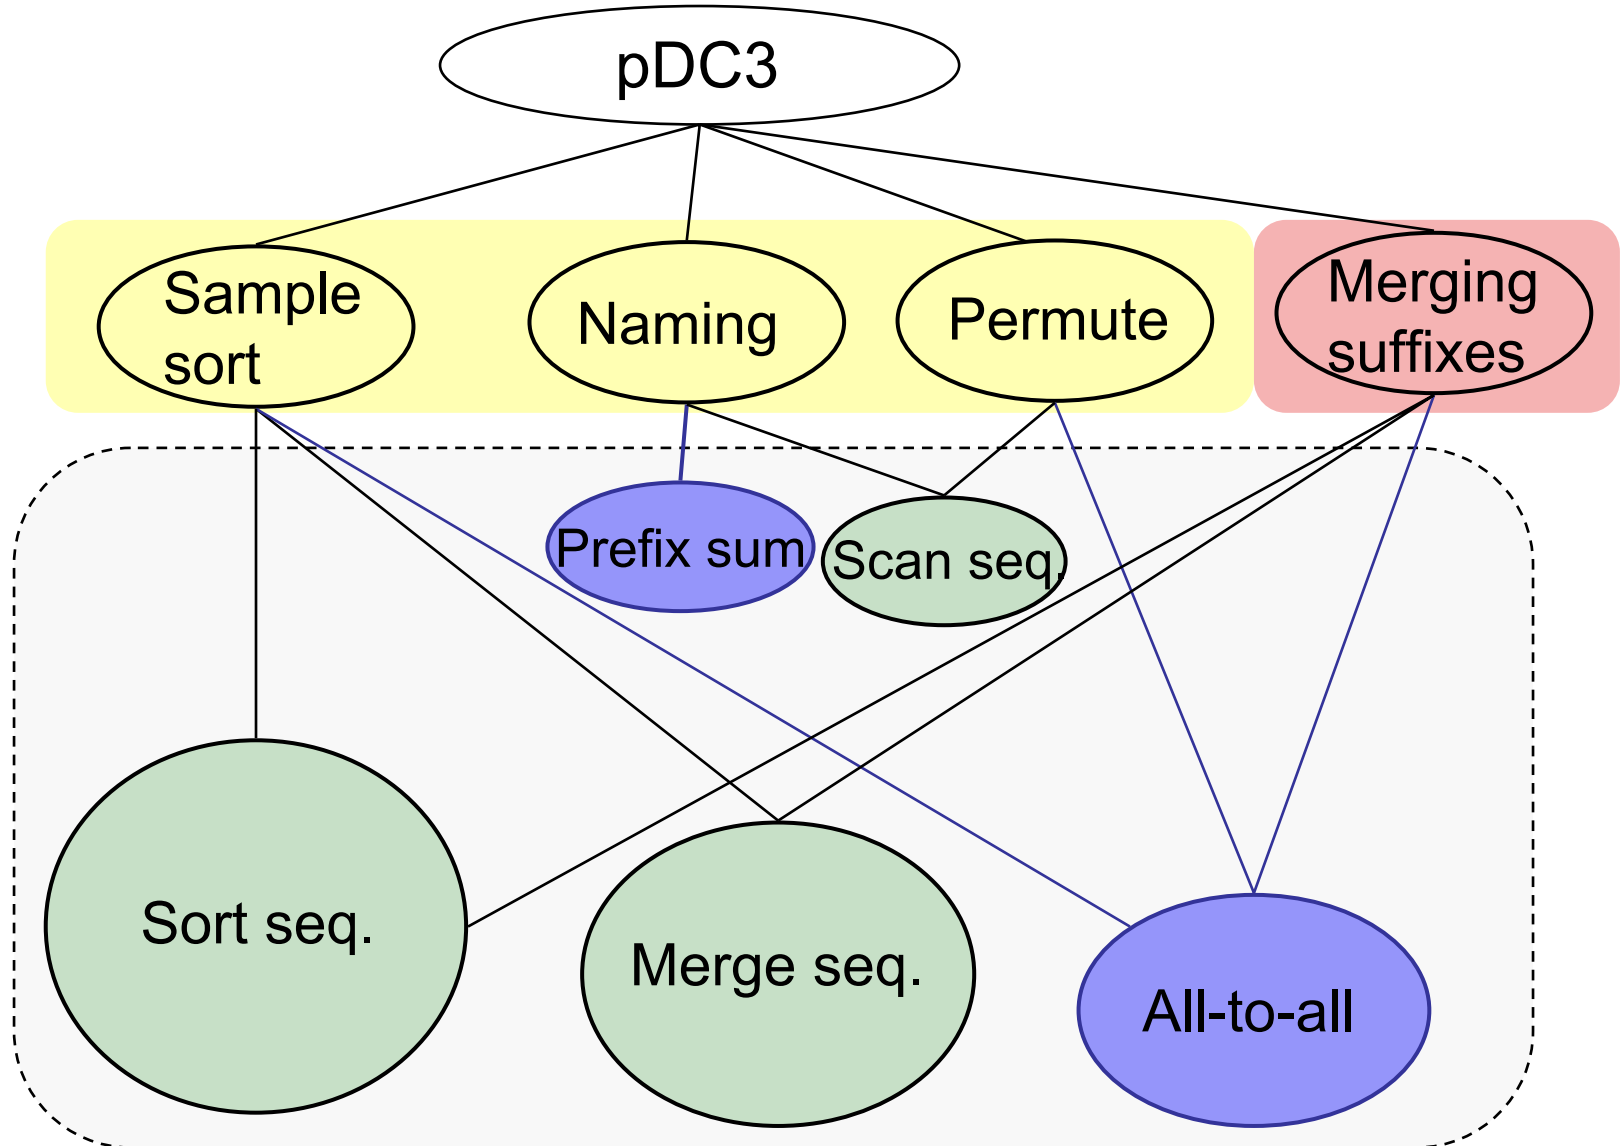
$ANA \leftrightarrow ANANA : (A,3) \leftrightarrow (A,4)$

$ANA \leftrightarrow NANA : (A,N,1) \leftrightarrow (N,A,3)$

Note: only parts of suffixes (triples/pairs) are compared

Overview Implementation

Detail
↓





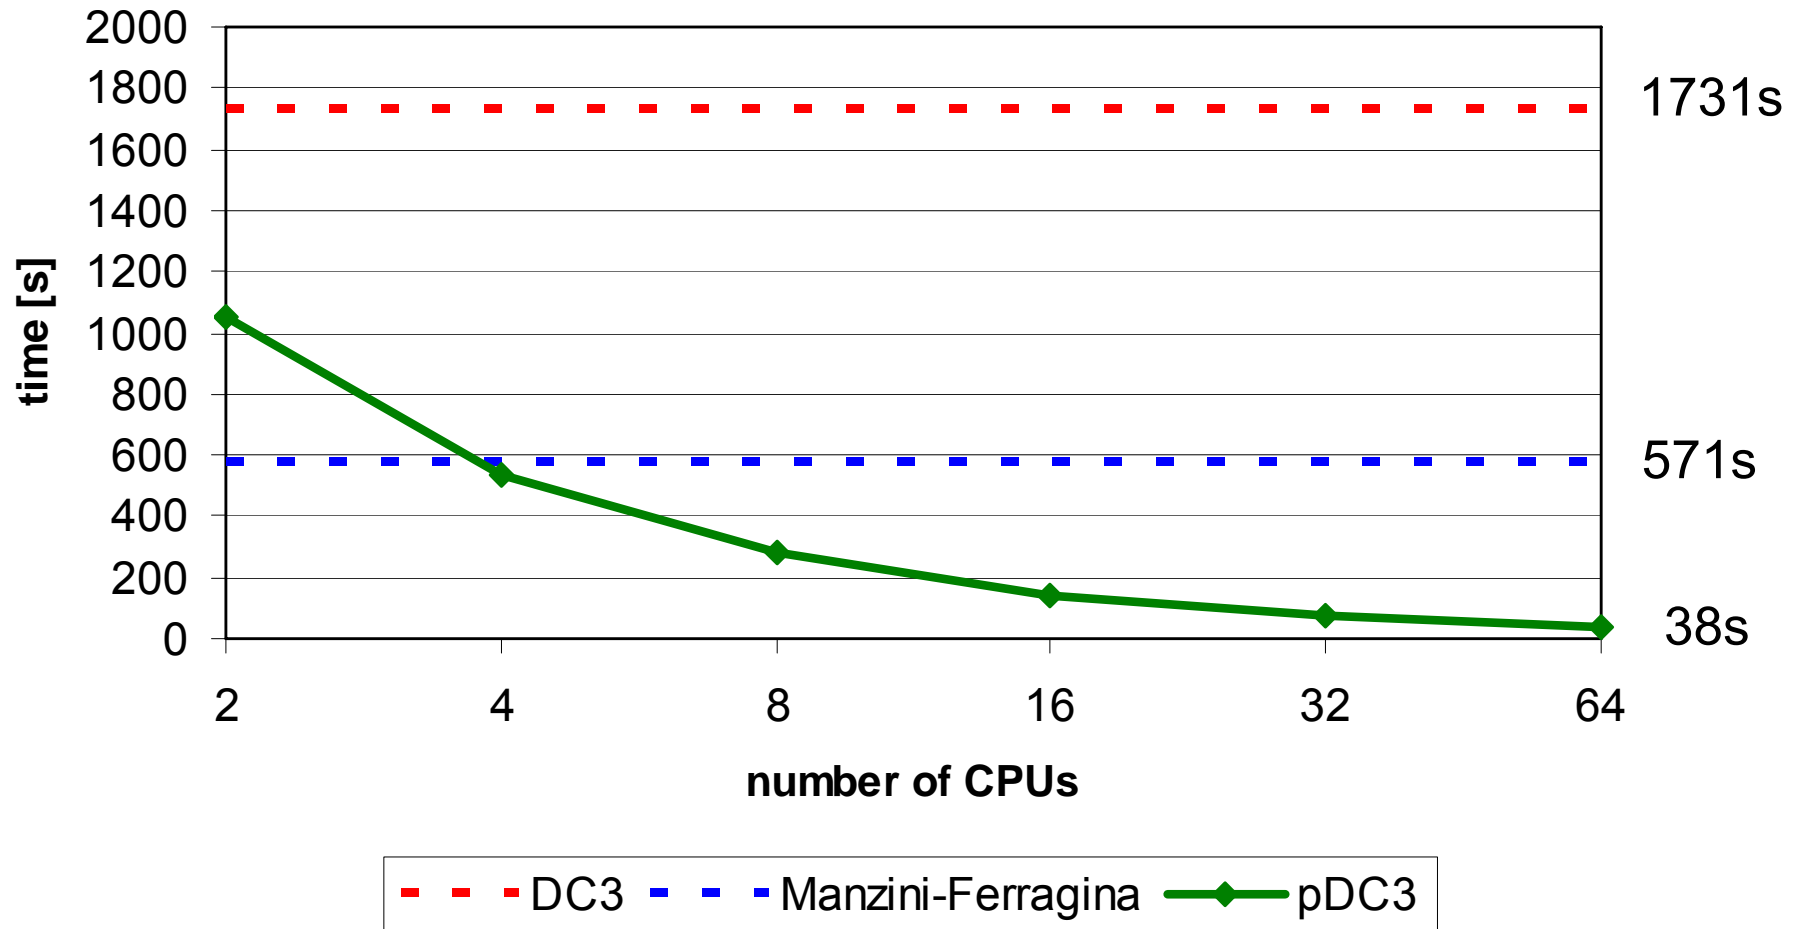
Karlsruhe



- Itanium cluster - University Karlsruhe (XC-Cluster)
(64 dual processor nodes available;
each node: 2x 1.5 GHz Itanium 2 and 12 GB main memory)
 - Linux source code (540 MB)
 - Human genome (~3 GB)
 - Gutenberg project (~3.2 GB)




Linux source code (540 MB)







Human genome (~3 GB)



Human genome



Gutenberg project (~3.2 GB)



- misc
- all-to-all
- merge max
- ▨ merge min
- sort max
- ▨ sort min




Jülich



- p690 IBM eServer Cluster1600 - FZ Jülich (JUMP-Cluster)
(16 SMP nodes available;
each node: 32x 1.7 GHz Power4+ and 128 GB main memory)
 - Human genome (~3 GB)

Human genome



Summary



- Suffix arrays:
lexicographically sorted array of the suffixes of a string
- Sample sort
Parallel sorting algorithm
- Overview pDC3
Sorting tuples and recursion
- Evaluation:
 - Influence of different sequential sorting algorithms
 - Good scalability
 - **Suffix array construction of human genome: 1.5 minutes**





Thank you for your
attention!
Questions?



Pseudo code of pDC3

Function $pDC3(T)$

$S := \langle (T[i, i+2]), i) : i \in [0, n), i \bmod 3 \neq 0 \rangle$

sort S by the first component

$P := \text{name}(S)$

if *the names in P are not unique* **then**

permute the $(r, i) \in P$ such that they are sorted by $(i \bmod 3, i \text{ div } 3)$

$SA^{12} := pDC3(\langle c : (c, i) \in P \rangle)$

$P := \langle (j+1, SA^{12}[j]) : j \in [0, 2n/3) \rangle$

permute P such that it is sorted by the second component

$S_0 := \langle (T[i], T[i+1], c', c'', i) : i \bmod 3 = 0, (c', i+1), (c'', i+2) \in P \rangle$

$S_1 := \langle (c, T[i], c', i) : i \bmod 3 = 1, (c, i), (c', i+1) \in P \rangle$

$S_2 := \langle (c, T[i], T[i+1], c'', i) : i \bmod 3 = 2, (c, i), (c'', i+2) \in P \rangle$

$S := \text{sort } S_0 \cup S_1 \cup S_2 \text{ using comparison function:}$

$(c, \dots) \in S_1 \cup S_2 \leq (d, \dots) \in S_1 \cup S_2 \Leftrightarrow c \leq d$

$(t, t', c', c'', i) \in S_0 \leq (u, u', d', d'', j) \in S_0 \Leftrightarrow (t, c') \leq (u, d')$

$(t, t', c', c'', i) \in S_0 \leq (d, u, d', j) \in S_1 \Leftrightarrow (t, c') \leq (u, d')$

$(t, t', c', c'', i) \in S_0 \leq (d, u, u', d'', j) \in S_2 \Leftrightarrow (t, t', c'') \leq (u, u', d'')$

return $\langle \text{last component of } s : s \in S \rangle$

1

2

3

4

5

6

7

8

9

10

11

12

13



Difference Cover



A **difference cover D modulo v** is a subset of $[0, v)$ such that for all $i \in [0, v)$ there exist $j, k \in D$ with $i \equiv k - j \pmod{v}$.

Example:

$\{1, 2\}$ is a difference cover modulo 3.

$\{1, 2, 4\}$ is a difference cover modulo 7.